

Agent Title Review

A practical review for consequential AI workflow actions

Candidate artifact by Russell Dudek for ServiceLink

Does the agent have clear title to act?

Use this review before approving a consequential agent behavior or when a change in model, data, policy, channel or workflow could alter the action boundary. It is not a compliance checklist; it is a cross-functional design conversation.

1. Intent

What is the user trying to accomplish? Which context is required? What ambiguity must be clarified? How will the agent confirm that it understood the request?

2. Evidence

Which sources materially support the action? Are they current, authenticated and consistent? What conflict or missing information must be shown rather than silently resolved?

3. Authority

May the agent complete, draft, recommend, confirm, reverse or never perform this action? Who owns the final decision? Is the action reversible and within a bounded scope?

4. Encumbrance

Which condition changes the rules: identity ambiguity, document conflict, regulated change, borrower distress, dispute, unusual risk, policy exception or channel limitation?

5. Recovery

How does the workflow clarify, offer alternatives, preserve context, warm-transfer, reverse safely and resume without forcing the user to start over?

6. Receipt

What record remains of intent, evidence, authority, exception, human decision, action and next step? Can a user understand it? Can the organization learn from it?

DISPOSITION

Clear to act	Controlled assist	Human review	Stop / redesign
Bounded authority, sufficient evidence, visible recovery.	Agent prepares or recommends; human retains consequential decision.	Exception or ambiguity requires a named human owner and evidence-rich handoff.	Authority, evidence or recovery is too weak for responsible deployment.

EVIDENCE TO COLLECT AFTER RELEASE

Task completion; user comprehension; overrides; handoff delay; recovery success; unsupported actions; recurring exceptions; workarounds; cycle time; rework; frontline feedback; time from observed failure to updated pattern.

Core design rule: confidence is not authority. A plausible answer does not establish the agent's right to act.