

Learn the work. Make behavior visible. Prove one reusable pattern.

This plan is a hypothesis for discovery, not a claim about ServiceLink's undisclosed priorities or internal architecture. The first priority is to understand the transaction, the users and the organizational decision rights before prescribing a solution.

DAYS 1-30 - READ THE TRANSACTION

People Build the stakeholder and user map.

Interview product, design, engineering, data science, operations, legal/compliance, client-facing teams and frontline users. Identify sponsors, decision owners, policy owners and the people absorbing exception work today.

Workflow Select one priority journey.

Map the current state from trigger to outcome across workflow UI, chat, SMS or voice where relevant. Capture user goals, system actions, handoffs, exception classes, recovery paths and evidence sources.

Behavior Inventory the agent's actual and intended authority.

Separate completion, recommendation, drafting, confirmation, escalation and prohibited actions. Identify where the same behavior is expressed differently across channels.

Evidence Establish a baseline.

Agree on existing qualitative and quantitative signals: cycle time, touches, rework, abandonment, overrides, escalation delay, missing information, user confidence, training burden and recurring failure modes.

Practice Audit the AI-enabled design workflow.

Understand how the team uses Figma, code-based prototypes, AI tools, research, design reviews and design-to-development handoff. Identify one low-risk acceleration opportunity.

30-DAY DELIVERABLES

- Current-state journey and exception map for one priority workflow.
- Stakeholder and decision-rights map.
- Agent behavior inventory across intent, evidence, authority, exception, recovery and receipt.
- Baseline measurement plan and prioritized research questions.
- Shared design principles for clarity, confirmation, escalation and recovery.

Decision at Day 30: Is the selected workflow bounded enough to prototype responsibly, important enough to matter and instrumentable enough to learn from?

DAYS 31-60 - PROTOTYPE THE CONTROL POINTS

Concept **Prototype the full behavior.**

Design clarification, confirmation, evidence, approval, exception, handoff, reversal and resume states.

Interaction **Test across surfaces.**

Keep workflow UI, chat, SMS and voice aligned to one intent and decision record.

Research **Use realistic exceptions.**

Test ambiguity, conflicting evidence, policy boundaries, distress and incomplete context.

Evaluation **Create an evaluation rubric.**

Assess success, comprehension, authority, escalation, recovery, operating impact and adoption.

Delivery **Pair with engineering and data science.**

Define deterministic rules, model inference, evidence, logging and implementation checks.

DAYS 61-90 - SHIP ONE REUSABLE PATTERN

Pilot **Deliver a bounded improvement.**

Prefer a narrow, observable slice with clear human authority.

Measure **Instrument outcomes.**

Monitor use, completion, cycle time, rework, overrides, exceptions, recovery and confidence.

Learn **Review failure modes.**

Convert confusion, unsupported action, poor handoff and workarounds into updated patterns.

Transfer **Package the reusable system.**

Publish states, copy, authority rules, accessibility, evidence, tests and implementation notes.

Roadmap **Recommend the next move.**

Identify where the pattern transfers, varies, or should be held.

90-DAY REVIEW

Ask: Did the workflow improve without hiding accountability? Can users understand and recover from agent action? Did decision effort decline or merely move? Which pattern is reusable across an adjacent EXOS experience? What improved in the team's AI design workflow? **Outcome:** one credible pattern, one evidence-backed reuse decision, and a clearer AI design operating system.