

Learn the work. Make behavior visible. Prove one reusable pattern.

This plan is a hypothesis for discovery, not a claim about ServiceLink's undisclosed priorities or internal architecture. The first obligation is to understand the transaction, the users, the operating constraints and the organizational decision rights before prescribing a solution.

DAYS 1-30 - READ THE TRANSACTION

People	Build the stakeholder and user map. Interview product, design, engineering, data science, operations, legal/compliance, client-facing teams and frontline users. Identify sponsors, decision owners, policy owners and the people absorbing exception work today.
Workflow	Select one priority journey. Map the current state from trigger to outcome across workflow UI, chat, SMS or voice where relevant. Capture user goals, system actions, handoffs, exception classes, recovery paths and evidence sources.
Behavior	Inventory actual and intended agent authority. Separate completion, recommendation, drafting, confirmation, escalation and prohibited actions. Identify where the same behavior is expressed differently across channels.
Evidence	Establish a baseline. Agree on existing qualitative and quantitative signals: cycle time, touches, rework, abandonment, overrides, escalation delay, missing information, user confidence, training burden and recurring failure modes.
Practice	Audit the AI-enabled design workflow. Understand how the team uses Figma, code-based prototypes, AI tools, research, design reviews and design-to-development handoff. Identify one low-risk acceleration opportunity.

30-DAY DELIVERABLES

- Current-state journey and exception map for one priority workflow.
- Stakeholder, policy-owner and decision-rights map.
- Agent behavior inventory across intent, evidence, authority, encumbrance, recovery and receipt.
- Baseline measurement plan and prioritized research questions.
- Shared design principles for clarity, confirmation, escalation and recovery.

Decision at Day 30: Is the selected workflow bounded enough to prototype responsibly, important enough to matter and instrumentable enough to learn from?

Prototype the full behavior, not only the happy path.

The middle phase turns discovery into a testable interaction system. The work should expose how the experience behaves when evidence conflicts, authority changes or a handoff fails—not merely demonstrate a polished normal flow.

DAYS 31-60 - DESIGN AND TEST

Concept	Prototype clarification through recovery. Design the intent check, evidence view, confirmation, approval, exception, handoff, reversal and resume states as one coherent experience.
Interaction	Keep every surface aligned. Ensure workflow UI, chat, SMS and voice reference one transaction state, one authority model and one decision record rather than creating channel-specific truths.
Research	Test realistic exceptions. Use ambiguity, conflicting evidence, policy boundaries, borrower distress, incomplete context and failed handoffs to expose where the design loses trust.
Evaluation	Create an AI interaction rubric. Assess task success, comprehension, evidence legibility, authority recognition, escalation quality, recovery, operating impact and adoption.
Delivery	Pair with engineering and data science. Define which behavior is deterministic, which depends on model inference, what evidence must be visible, what gets logged and which implementation checks block release.

60-DAY PROTOTYPE PACKAGE

- Testable end-to-end prototype covering normal, ambiguous and exception states.
- Conversation patterns for clarification, confirmation, refusal, escalation and warm transfer.
- Agent Title Review specifying intent, evidence, authority, encumbrance, recovery and receipt.
- Evaluation rubric with usability, operating and risk signals.
- Implementation annotations for state, data, copy, accessibility, analytics and quality checks.

Decision at Day 60: Does the concept improve the user's ability to act and recover without hiding authority, moving work downstream or creating an unreviewable decision?

Deliver one credible pattern and a better next decision.

The final phase should produce more than a demonstration. It should create implementation evidence, a measurable operating agreement and a reusable pattern that improves the next product decision.

DAYS 61-90 - DELIVER AND LEARN

Pilot	Ship a bounded workflow improvement. Prefer a narrow, observable slice with clear human authority and a recovery path over a broad demonstration that cannot be trusted operationally.
Measure	Instrument the experience. Monitor use, completion, cycle time, rework, overrides, exception frequency, escalation delay, recovery success and user confidence.
Learn	Review failure modes weekly. Convert confusion, unsupported action, poor handoff, repeated clarification and workarounds into revised patterns and product hypotheses.
Transfer	Package the reusable system. Publish states, copy, authority rules, accessibility guidance, evidence requirements, evaluation cases and implementation notes.
Roadmap	Recommend the next move. Identify where the pattern transfers directly, where legitimate workflow variation requires adaptation and where the organization should hold rather than automate.

90-DAY REVIEW SCORECARD

Dimension	Review question	Evidence to bring
User and workflow	Did the experience reduce effort, delay or avoidable rework?	Task outcomes, cycle time, touches, usability findings and frontline feedback.
Trust and authority	Could users explain what the agent knew, what it could do and when a person remained accountable?	Comprehension tests, approval/override behavior, escalation quality and trust-receipt review.
Recovery and risk	Did exceptions return to a good state without lost context, silent failure or policy ambiguity?	Exception cases, recovery success, failure analysis and legal/compliance feedback.
Reuse and practice	Is the pattern reusable, and did the team's AI design workflow become faster or more rigorous?	Pattern adoption, implementation effort, review quality and the next prioritized workflow.

REVIEW CONVERSATION

Ask: Did decision effort decline or merely move? Which assumption proved wrong? Where did human judgment add the most value? What should become a reusable design-system pattern? Which adjacent EXOS experience now has enough evidence to explore, and which should remain intentionally unautomated?

Outcome by Day 90: one shipped or implementation-ready bounded improvement, one evidence-backed reuse decision, one visible hard-objection or failure-mode record, and a clearer operating system for designing trustworthy agent behavior.